1. Introduction to AI Agents

An AI agent is a software system that uses a Large Language Model (LLM) to understand an objective, reason about possible actions, interact with tools or external systems, and work toward completing a task. Unlike a basic question-answering application, an agent can decide which action to take, inspect the result, and continue when more work is needed.

1.1 Key Components of an AI Agent

- **LLM:** Provides language understanding, reasoning, and generation capabilities.
- **Instructions:** Define the agent's role, goals, constraints, and expected behavior.
- **Tools:** Allow the agent to call APIs, query databases, search the web, or execute code.
- **Memory and state:** Preserve relevant information across steps or interactions.
- **Planning:** Breaks an objective into smaller actions and determines their order.
- **Execution and feedback:** Runs actions, observes results, and decides whether to continue.

1.2 LLM vs. AI Agent

Feature	Traditional LLM application	AI agent
Input	Usually a prompt or message	An objective or task
Output	Primarily generated text	Text, decisions, or executed actions
Tool use	Often explicitly wired into the application	Can select tools based on the task and available permissions
Planning	Usually limited to the current generation	Can plan and iterate across multiple steps
State	Often limited to conversation context	Can maintain workflow state and persistent memory
Autonomy	Generally lower	Can perform multi-step work within defined boundaries

1.3 Basic Agent Execution Loop

```
User objective
      |
      v
LLM reasoning and planning
      |
      v
Select an action or tool
      |
      v
Execute the action
      |
      v
Observe the result
      |
      +----> Continue / replan if needed
      |
      v
Return the final response
```

2. LangChain and LangGraph

2.1 Introduction to LangChain

LangChain is an open-source framework for developing applications powered by language models. It provides abstractions and integrations for connecting models with prompts, tools, retrievers, vector stores, and external services.

Core components

- **Models:** Language models accessed through providers such as OpenAI, Anthropic, Google, OpenRouter, or local runtimes.
- **Prompts:** Reusable instructions and message templates supplied to a model.
- **Tools:** Functions that let an agent interact with APIs, databases, search systems, and other services.
- **Retrievers:** Components that fetch relevant information from document collections or knowledge bases.
- **Output parsers:** Convert model output into structured formats such as JSON or typed objects.
- **Agent execution:** Combines model reasoning and tool invocation into an iterative workflow.

Example: defining a prompt and a tool

```
from langchain_core.prompts import ChatPromptTemplate
from langchain_core.tools import tool

prompt = ChatPromptTemplate.from_messages([
    ("system", "You are a helpful AI assistant."),
    ("human", "{question}")
])

@tool
def calculator(a: int, b: int) -> int:
    """Add two integers."""
    return a + b
```

2.2 Introduction to LangGraph

LangGraph is a framework for building stateful, controllable, and potentially long-running agent workflows. It models an application as a graph made up of nodes, edges, and shared state. It is particularly useful when developers need explicit control over routing, cycles, persistence, and human approval.

Core concepts of LangGraph

- **State:** The shared data passed between graph steps, such as the user request, messages, tool results, and status.
- **Nodes:** Processing units that can call an LLM, execute a function, invoke a tool, or transform data.
- **Edges:** Define transitions between nodes. Conditional edges choose the next step based on the current state.
- **Checkpointing:** Saves graph state so an execution can be resumed or inspected.
- **Human-in-the-loop:** Allows execution to pause for review or approval before continuing.

Example: defining graph state

```
from typing_extensions import TypedDict

class AgentState(TypedDict):
    user_query: str
    response: str
    tool_result: str
```

2.3 LangChain vs. LangGraph

Dimension	LangChain	LangGraph
Primary focus	LLM application building blocks and agent abstractions	Orchestration of stateful workflows and agents
Workflow model	Chains and agent abstractions	Graphs composed of nodes and edges
Control flow	Often managed through higher-level abstractions	Explicit transitions and conditional routing
Cycles	Available through agent execution patterns	Naturally represented as graph cycles
State and persistence	Integrates with state and persistence options	Explicit graph state and checkpointing
Best suited for	Model, prompt, tool, and retrieval integration	Complex workflows, routing, and long-running execution

Practical rule of thumb: use LangChain abstractions when you need convenient model and tool integration; use LangGraph when you need explicit workflow control, state transitions, cycles, persistence, or approval steps.

2.4 Illustrative LangGraph Workflow

```
START
  |
  v
Input processing
  |
  v
LLM agent
  |
  v
Decision: tool required?
  |                 |
 Yes                No
  |                 |
  v                 v
Call tool       Final answer
  |
  v
Update state
  |
  +------> Return to LLM agent
```

3. Agent Lifecycle

The agent lifecycle describes the stages an agent follows from initialization and task understanding through planning, action, observation, and completion. Monitoring and evaluation help determine whether the agent is reliable and effective.

3.1 Lifecycle Stages

Stage 1 — Initialization

The system configures the model, system instructions, available tools, permissions, memory, and execution limits.

```
agent_config = {
    "name": "ResearchAgent",
    "role": "Research Assistant",
    "goal": "Find reliable information",
    "tools": ["web_search", "document_search"]
}
```

Stage 2 — Perception and input processing

The agent interprets the request, identifies the user's intent, extracts relevant entities and constraints, and determines the expected result. For example, “Find recent research papers on multi-agent systems” implies a research task, a topic, a freshness requirement, and a list or summary of papers as the expected output.

Stage 3 — Planning and reasoning

The agent decomposes the objective into subtasks, identifies dependencies, selects tools, and determines an execution strategy. A research task might require searching academic sources, filtering papers, extracting findings, and preparing a summary.

Stage 4 — Tool selection and execution

The agent chooses and invokes tools that are appropriate for the task. Common tools include web search, database queries, Python execution, file retrieval, and external APIs. Tool calls should be validated and restricted by permissions.

Stage 5 — Observation

The agent receives tool results and checks whether execution succeeded, whether the returned information is relevant, and whether the objective has been met. Tool output becomes new context for the next reasoning step.

Stage 6 — Iteration and replanning

If the result is incomplete or incorrect, the agent can revise its plan and try another action. For example, it may refine a search query after receiving irrelevant results. Iteration should be bounded by time, cost, and maximum-step limits.

Stage 7 — Completion and termination

The agent stops when the objective is achieved, a defined stopping condition is reached, the maximum iteration limit is exceeded, or the task cannot be completed safely or successfully. It then returns an appropriate result or explanation.

Stage 8 — Monitoring and evaluation

The system measures accuracy, task success rate, tool-selection quality, latency, token usage, cost, and failure rate. Traces and logs help developers diagnose errors and improve the workflow.

3.2 Complete Agent Lifecycle

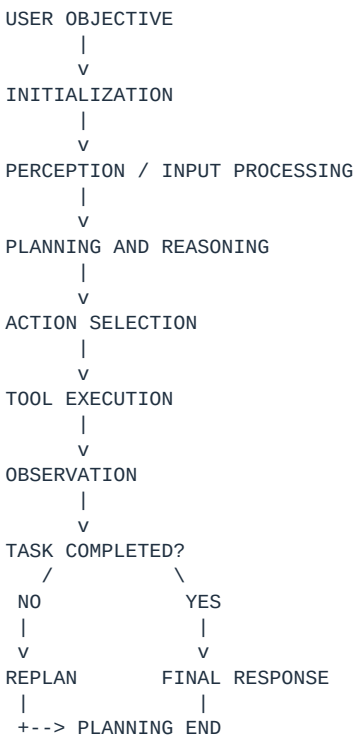

4. Multi-Agent Coordination Patterns

4.1 What is a Multi-Agent System?

A multi-agent system consists of multiple autonomous or semi-autonomous agents that cooperate, coordinate, or delegate work to achieve an objective. Each agent may have a specialized role, instructions, tools, context, and responsibilities. A coordinator or communication mechanism may manage how their work is combined.

Example: a research system may use a Research Agent to collect sources, an Analysis Agent to identify insights, a Verification Agent to check claims, and a Writing Agent to prepare the final report.

4.2 Major Coordination Patterns

Pattern 1 — Supervisor

A supervisor receives the request, decomposes it, delegates subtasks to specialized agents, reviews their results, and produces or requests the final output.

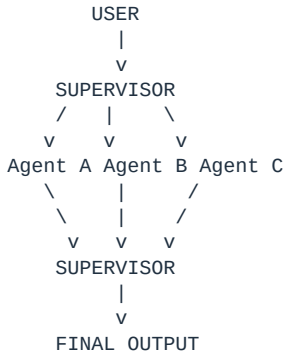

Typical uses: complex research, customer support, business intelligence, and workflow automation.

Pattern 2 — Sequential

Agents execute in a predefined order. The output from one stage becomes the input to the next stage.

```
User request -> Research -> Analysis -> Verification -> Writing -> Final output
```

Typical uses: document processing, report generation, and ordered data-analysis pipelines.

Pattern 3 — Parallel

Independent agents work simultaneously on separate subtasks. An aggregator combines their results after execution.

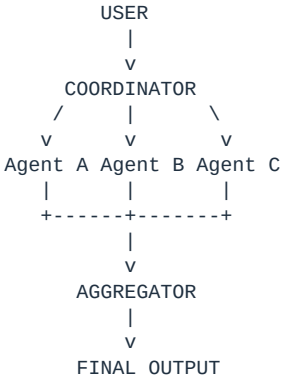

Typical uses: market research, multi-source analysis, and parallel data extraction.

Pattern 4 — Hierarchical

Agents are arranged in multiple levels. A high-level agent delegates work to team leads or sub-supervisors, who coordinate specialized agents.

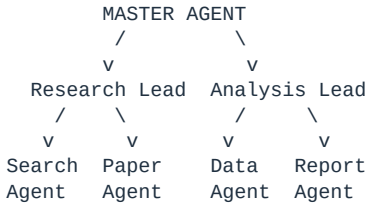

Typical uses: large-scale research and complex enterprise processes.

Pattern 5 — Network / decentralized

Agents communicate directly with other agents and may decide which agent should handle the next step. There is no requirement for one central supervisor to direct every transition.

Typical uses: collaborative problem-solving and distributed decision-making. This pattern requires careful control of routing, loops, and shared context.

Pattern 6 — Handoff

An agent transfers control to another agent when the task requires a different specialization. The receiving agent takes responsibility for the next part of the interaction.

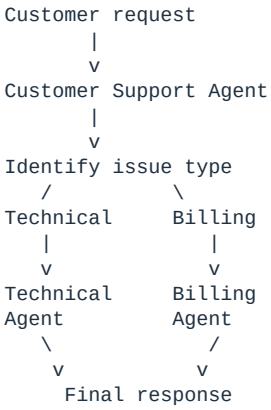

Typical uses: customer support, technical troubleshooting, and service workflows.

Pattern 7 — Reflection and critique

A generator agent creates an initial output, while a critic or evaluator agent reviews it against defined criteria. Feedback is used to revise the output until it is accepted or a stopping condition is reached.

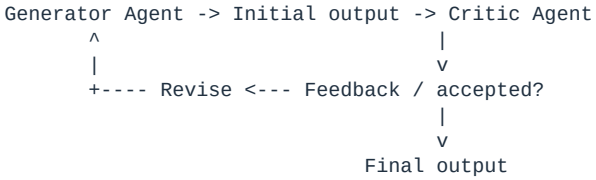

Typical uses: code generation, report validation, and quality improvement.

4.3 Pattern Comparison

Pattern	Coordination style	Typical use
Supervisor	Centralized delegation and review	Complex task delegation
Sequential	Fixed step-by-step flow	Ordered workflows
Parallel	Concurrent independent work	Multi-source research
Hierarchical	Multiple levels of delegation	Enterprise-scale tasks
Network	Decentralized agent interaction	Collaborative problem-solving
Handoff	Dynamic transfer of control	Customer support
Reflection	Generation followed by critique	Quality improvement

4.4 Common Challenges

- Communication overhead and increased latency.
- Higher LLM token consumption and operating cost.
- Conflicting outputs or inconsistent assumptions between agents.
- Shared-state consistency and context management.
- Infinite loops or repeated delegation.
- Tool failures, retries, and partial execution.
- Security, authorization, and least-privilege access.

A multi-agent design is not automatically better than a single-agent design. Use multiple agents when specialization, parallelism, or independent review provides a measurable benefit.

5. Model Context Protocol (MCP)

5.1 What is MCP?

The Model Context Protocol (MCP) is an open protocol that standardizes how AI applications connect to external tools, data sources, and services. It provides a common interface so an AI application can discover and use capabilities exposed by MCP servers, rather than requiring a separate custom integration for every service.

5.2 MCP Architecture

MCP uses a Host–Client–Server architecture:

- **MCP Host:** The AI application that manages the user interaction and coordinates MCP connections.
- **MCP Client:** A component within the host that communicates with an MCP server.
- **MCP Server:** Exposes capabilities such as tools, resources, and prompts to an MCP client.

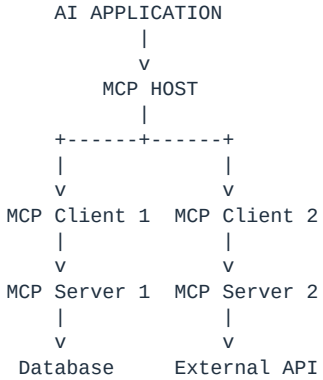

Each MCP client communicates with an MCP server. A host may manage connections to multiple servers.

5.3 Core MCP Primitives

Tools

Tools are executable capabilities exposed by an MCP server. Examples include searching a database, creating a file, retrieving customer information, or invoking an external operation. Tool execution should be subject to validation and authorization.

Resources

Resources expose contextual data that an application can read, such as documents, database schemas, configuration files, or project information.

Prompts

Prompts are reusable prompt templates that an MCP server can make available to a client, such as templates for code review, research, or data analysis.

5.4 MCP Communication Flow

```
AI APPLICATION
      |
      v
MCP CLIENT
      |
      v
Connect to MCP server and discover capabilities
      |
      v
MCP SERVER exposes tools, resources, and prompts
      |
      v
Application selects an appropriate capability
      |
      v
Request is sent to the server
      |
      v
Server returns a result
      |
      v
AI APPLICATION uses the result
```

MCP uses JSON-RPC-based messages. Common transports include STDIO for local integrations and Streamable HTTP for networked integrations.

5.5 MCP vs. Multi-Agent Coordination

Dimension	Multi-agent coordination	Model Context Protocol
Purpose	Coordinates agents working on a task	Connects AI applications to external capabilities
Main focus	Collaboration and task delegation	Interoperability and context exchange
Components	Agents, supervisors, orchestrators	Hosts, clients, and servers
Communication	Agent-to-agent or orchestrator-mediated	Client-server protocol
Example	Research Agent delegates analysis to another agent	An agent queries a database through an MCP server

Important distinction: MCP does not define how multiple agents should collaborate. An application can use LangGraph to coordinate agents and MCP to connect those agents to external tools and data sources.

6. End-to-End Practical Implementation

6.1 Project: Multi-Agent Research Assistant

Objective: Build a research assistant that accepts a user query, gathers relevant information, analyzes and verifies the findings, and generates a structured research report.

6.2 Proposed Agents

- **Supervisor Agent:** Manages the overall workflow and delegates tasks.
- **Research Agent:** Collects relevant information from approved sources.
- **Analysis Agent:** Extracts themes, comparisons, and key insights.
- **Verification Agent:** Checks whether claims are supported by the available evidence.
- **Report Agent:** Produces the final structured report.

6.3 High-Level Architecture

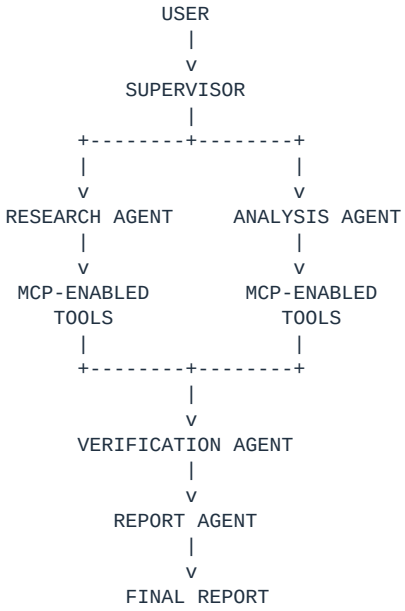

The diagram is illustrative. In a real implementation, the supervisor may choose sequential or parallel execution, and the verification stage may send work back for revision.

6.4 Suggested Technology Stack

Layer	Suggested technology	Role
Programming language	Python	Application implementation
LLM integration	LangChain	Model and tool integration
Orchestration	LangGraph	Agent state, routing, and workflow control
Model provider	OpenRouter or another provider	Access to hosted language models
External integration	MCP	Standardized access to tools and data
Vector database	ChromaDB or FAISS	Optional semantic retrieval
Backend	FastAPI	Expose application functionality through APIs
Frontend	Streamlit	Research assistant user interface
Monitoring	LangSmith or equivalent tracing	Inspect execution and evaluate behavior

6.5 Implementation Roadmap

- Configure the LLM provider and securely load credentials from environment variables.
- Define the responsibilities, system prompts, and output formats for each agent.
- Implement research and information-retrieval tools.
- Create a LangGraph state schema to hold the query, intermediate findings, verification status, and final report.
- Implement each agent as a graph node or a clearly defined subgraph.
- Define graph edges, conditional routing, and any parallel execution.
- Connect approved external capabilities through MCP clients and servers.
- Add timeouts, retries, maximum-step limits, and clear error handling.
- Add checkpointing and human approval where required.
- Test the workflow using representative research queries and evaluate quality, latency, and cost.

6.6 Simplified Workflow

```
1. User submits a research question.
2. Supervisor interprets the objective and creates subtasks.
3. Research Agent retrieves relevant material using approved tools.
4. Analysis Agent extracts findings and supporting evidence.
5. Verification Agent checks the findings.
6. If verification fails, the supervisor requests additional research or revision.
7. Report Agent prepares the final structured report.
8. The application returns the report to the user.
```

End of requested content — Sections 1 through 6.